

Experiment No.6

Title: Applying similarity measures on the textual datasets

Batch: B4

Roll No.: 16010420133

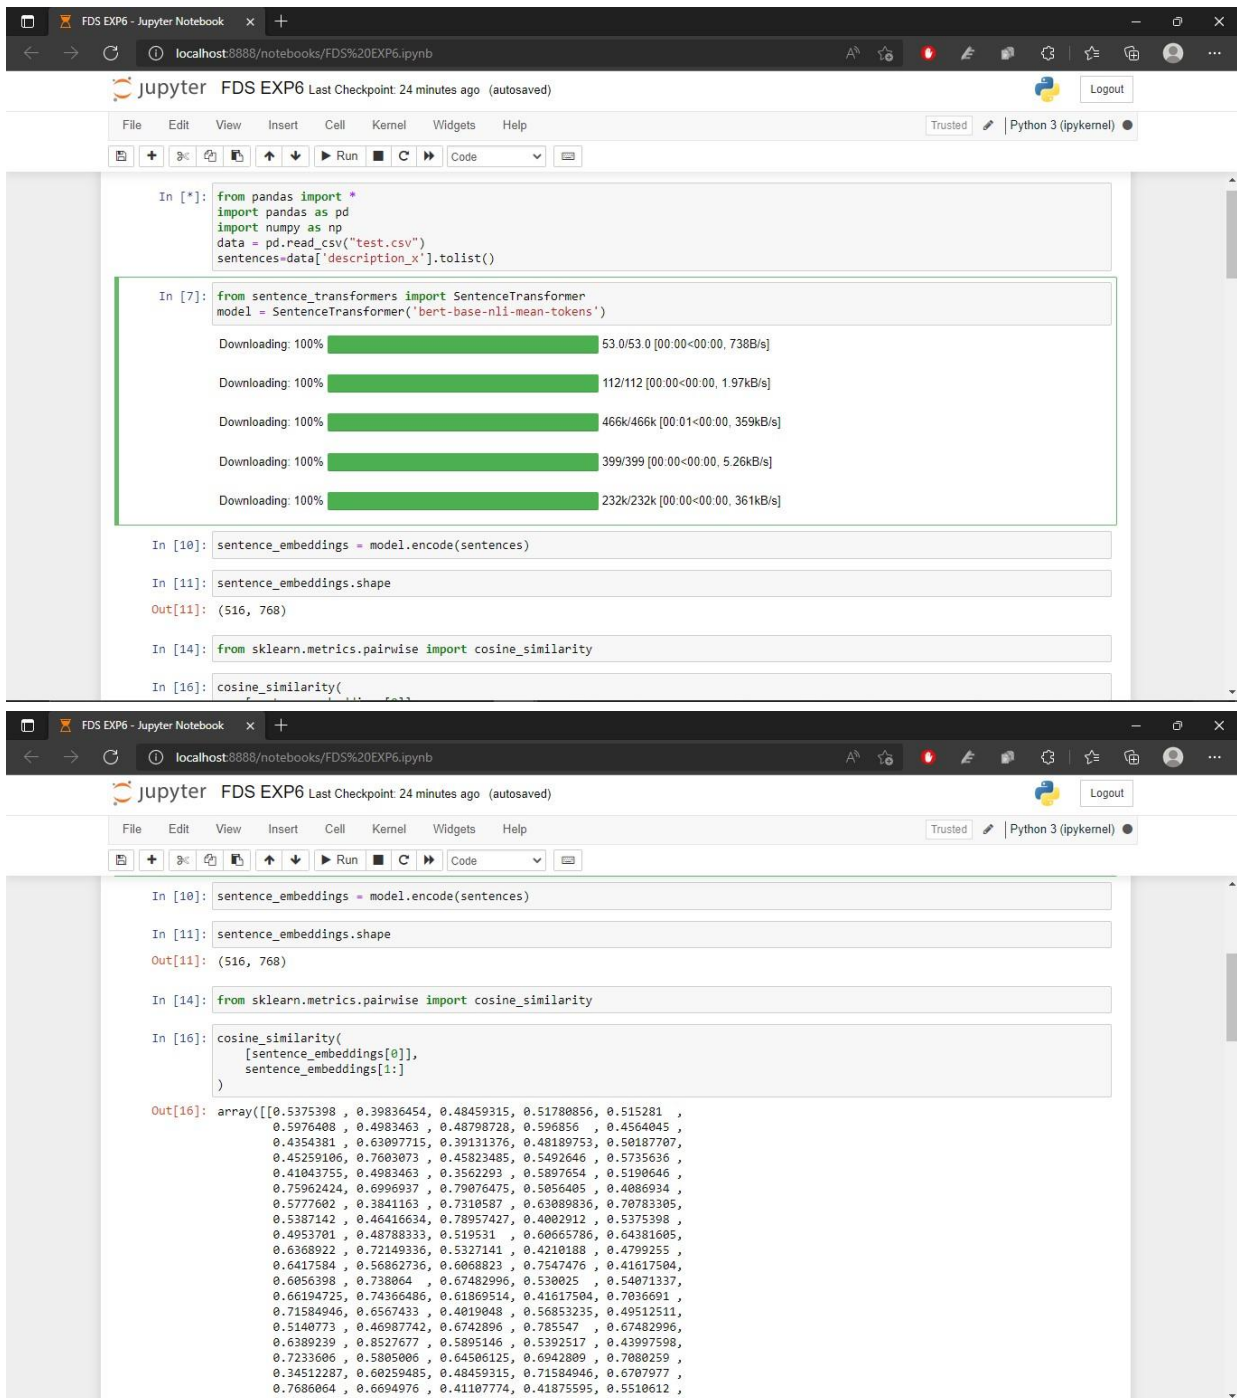
Experiment No.: 6

Aim: Applying similarity measures on the textual datasets using cosine distance and BERT.

Resources needed: Python

Procedure / Approach / Algorithm / Activity Diagram:

1. Convert some given set of statements into dense vectors and calculate the cosine distance between them using BERT model in python.



The first screenshot shows the initial steps of the Jupyter Notebook:

```
In [*]: from pandas import *
import pandas as pd
import numpy as np
data = pd.read_csv("test.csv")
sentences=data['description_x'].tolist()

In [7]: from sentence_transformers import SentenceTransformer
model = SentenceTransformer('bert-base-nli-mean-tokens')

Downloading: 100% ██████████ 53.0/53.0 [00:00<00:00, 738B/s]
Downloading: 100% ██████████ 112/112 [00:00<00:00, 1.97kB/s]
Downloading: 100% ██████████ 466k/466k [00:01<00:00, 359kB/s]
Downloading: 100% ██████████ 399/399 [00:00<00:00, 5.26kB/s]
Downloading: 100% ██████████ 232k/232k [00:00<00:00, 361kB/s]

In [10]: sentence_embeddings = model.encode(sentences)

In [11]: sentence_embeddings.shape
Out[11]: (516, 768)

In [14]: from sklearn.metrics.pairwise import cosine_similarity

In [16]: cosine_similarity(
```

The second screenshot shows the continuation of the code and the output of the cosine similarity calculation:

```
In [10]: sentence_embeddings = model.encode(sentences)

In [11]: sentence_embeddings.shape
Out[11]: (516, 768)

In [14]: from sklearn.metrics.pairwise import cosine_similarity

In [16]: cosine_similarity(
[sentence_embeddings[0]],
[sentence_embeddings[1:]]
)
Out[16]: array([[0.5375398, 0.39836454, 0.48459315, 0.51788856, 0.515281,
0.5976408, 0.4983463, 0.48798728, 0.596856, 0.4564045,
0.4354381, 0.63097715, 0.39131376, 0.48189753, 0.50187707,
0.45259106, 0.7603073, 0.45823485, 0.5492646, 0.5735636,
0.41043755, 0.4983463, 0.3562293, 0.5897654, 0.5190646,
0.75962424, 0.6996937, 0.79076475, 0.5056405, 0.4086934,
0.5777602, 0.3841163, 0.7310587, 0.63089836, 0.70783305,
0.5387142, 0.46416634, 0.78957427, 0.4002912, 0.5375398,
0.4953701, 0.48788333, 0.519531, 0.60665786, 0.64381605,
0.6368922, 0.72149336, 0.5327141, 0.4210188, 0.4799255,
0.6417584, 0.56862736, 0.6068823, 0.7547476, 0.41617584,
0.6056398, 0.738064, 0.67482996, 0.530025, 0.54071337,
0.66194725, 0.74366486, 0.61869514, 0.41617504, 0.7036691,
0.71584946, 0.6567433, 0.4019048, 0.56853235, 0.49512511,
0.5140773, 0.46987742, 0.6742896, 0.785547, 0.67482996,
0.6389239, 0.8527677, 0.5895146, 0.5392517, 0.43997598,
0.7233606, 0.5805006, 0.64506125, 0.6942809, 0.7080259,
0.34512287, 0.60259485, 0.48459315, 0.71584946, 0.6707977,
0.7686064, 0.6694976, 0.41107774, 0.41875595, 0.5510612,
0.436796, 0.7305416, 0.4979381, 0.6383847, 0.47304556])
```

```
In [16]: cosine_similarity(
          [sentence_embeddings[0]],
          sentence_embeddings[1:]
        )
```

```
Out[16]: array([[0.5375398 , 0.39836454, 0.48459315, 0.51780856, 0.515281 ,
                  0.5976408 , 0.4983463 , 0.48798728, 0.596856 , 0.4564045 ,
                  0.4354381 , 0.63097715, 0.39131376, 0.48189753, 0.50187707,
                  0.45259106, 0.7603073 , 0.45823485, 0.5492646 , 0.5735636 ,
                  0.41043755, 0.4983463 , 0.3562293 , 0.5897654 , 0.5190646 ,
                  0.75962424, 0.6996937 , 0.79076475, 0.5056405 , 0.4086934 ,
                  0.5777602 , 0.3841163 , 0.7310587 , 0.63089836, 0.70783305,
                  0.5387142 , 0.46416634, 0.78957427, 0.4002912 , 0.5375398 ,
                  0.4953701 , 0.48788333, 0.519531 , 0.60665786, 0.64381605,
                  0.6368922 , 0.72149336, 0.5327141 , 0.4210188 , 0.4799255 ,
                  0.6417584 , 0.56862736, 0.6068823 , 0.7547476 , 0.41617504,
                  0.6056398 , 0.738064 , 0.67482996, 0.530025 , 0.54071337,
                  0.66194725, 0.74366486, 0.61869514, 0.41617504, 0.7036691 ,
                  0.71584946, 0.6567433 , 0.4019048 , 0.56853235, 0.49512511,
                  0.5140773 , 0.46987742, 0.6742896 , 0.785547 , 0.67482996,
                  0.6389239 , 0.8527677 , 0.5895146 , 0.5392517 , 0.43997598,
                  0.7233606 , 0.5805006 , 0.64506125, 0.6942809 , 0.7080259 ,
                  0.34512287, 0.60259485, 0.48459315, 0.71584946, 0.6707977 ,
                  0.7686064 , 0.6694976 , 0.41107774, 0.41875595, 0.5510612 ,
                  0.626706 , 0.7295516 , 0.4070281 , 0.6282847 , 0.67294586,
                  0.721545 , 0.3230321 , 0.47742856, 0.60753655, 0.426233 ,
                  0.62356377, 0.70055604, 0.39307052, 0.47112888, 0.49258387,
                  0.5443704 , 0.4285535 , 0.6545238 , 0.6093745 , 0.47207683,
                  0.48654285, 0.5510993 , 0.38145378, 0.5266507 , 0.62236273,
                  0.5043524 , 0.4003529 , 0.70268047, 0.5691283 , 0.41011897,
                  0.5370861 , 0.39089447, 0.5189902 , 0.54176575, 0.56605315,
                  0.73110896, 0.37961233, 0.71584946, 0.56426513, 0.5912481 ,
                  0.5852668 , 0.63003564, 0.5758078 , 0.38299286, 0.48341605,
                  0.47427866, 0.5842754 , 0.6000510 , 0.36060637, 0.6250140 ])
```

Questions:

1. Define cosine similarity with an example.

Cosine similarity is one of the metric to measure the text-similarity between two documents irrespective of their size in Natural language Processing. A word is represented into a vector form. The text documents are represented in n-dimensional vector space.

Mathematically, Cosine similarity metric measures the cosine of the angle between two n-dimensional vectors projected in a multi-dimensional space. The Cosine similarity of two documents will range from 0 to 1. If the Cosine similarity score is 1, it means two vectors have the same orientation. The value closer to 0 indicates that the two documents have less similarity. The mathematical equation of Cosine similarity between two non-zero vectors is:

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}},$$

Eg:

doc_1 = "Data is the oil of the digital economy"
doc_2 = "Data is a new oil"

Vector representation of the document

doc_1_vector = [1, 1, 1, 1, 0, 1, 1, 2]

doc_2_vector = [1, 0, 0, 1, 1, 0, 1, 0]

	data	digital	economy	is	new	of	oil	the
doc_1	1	1	1	1	0	1	1	2
doc_2	1	0	0	1	1	0	1	0

$$A \cdot B = \sum_{i=1}^n A_i B_i$$

$$= (1 * 1) + (1 * 0) + (1 * 0) + (1 * 1) + (0 * 1) + (1 * 0) + (1 * 1) + (2 * 0) \\ = 3$$

$$\sqrt{\sum_{i=1}^n A_i^2} = \sqrt{1+1+1+1+0+1+1+4} = \sqrt{10}$$

$$\sqrt{\sum_{i=1}^n B_i^2} = \sqrt{1+0+0+1+1+0+1+0} = \sqrt{4}$$

$$\text{cosine similarity} = \cos\theta = \frac{A \cdot B}{|A||B|} = \frac{3}{\sqrt{10} * \sqrt{4}} = 0.4743$$

2. What are similarity measures applicable to textual texts other than cosine distance.

List and explain at least 3 of them.

- 1) edit distance
- 2) cosine distance
- 3) Jaro distance
- 4) n-Gram distance
- 5) longest common subsequence

Edit Distance:

Levenshtein distance is the most frequently used algorithm. It was founded by the Russian scientist, Vladimir Levenshtein to calculate the similarities between two strings. This is also known as the Edit distance-based algorithm as it computes the number of edits required to transform one string to another. The edits count the following as an operation:

- Insertion of a character
- Deletion of a character
- Substitution of a character

More the number of operations, less is the similarity between the two strings.

Jaro Distance

Jaro Similarity is the measure of similarity between two strings. The value of Jaro distance ranges from 0 to 1, where 1 means the strings are equal and 0 means no similarity between the two strings.

The Jaro Similarity is calculated using the following formula

$$Jaro\ similarity = \begin{cases} 0, & \text{if } m=0 \\ \frac{1}{3} \left(\frac{m}{|s1|} + \frac{m}{|s2|} + \frac{m-t}{m} \right), & \text{for } m \neq 0 \end{cases}$$

where:

- m is the number of matching characters
- t is half the number of transpositions
- where |s1| and |s2| are the lengths of strings s1 and s2 respectively.

n-Gram distance

The main idea behind n-gram similarity is generalizing the concept of the longest common subsequence to encompass n-grams, rather than just unigrams. We formulate n-gram similarity as a function s_n , where n is a fixed parameter. s_1 is equivalent to the unigram similarity function s.

Outcomes:

CO 2 Comprehend descriptive and proximity measures of data

Conclusion:

I have implemented the cosine distance on the text stored in the textual dataset using BERT in python.

Grade: AA / AB / BB / BC / CC / CD /DD

Signature of faculty in-charge with date

References:

Books/ Journals/ Websites:

1. Han, Kamber, "Data Mining Concepts and Techniques", Morgan Kaufmann 3rd Edition
2. Tan, Pang-Ning, Michael Steinbach, and Vipin Kumar. Introduction to data mining. Pearson Education India, 2016.